



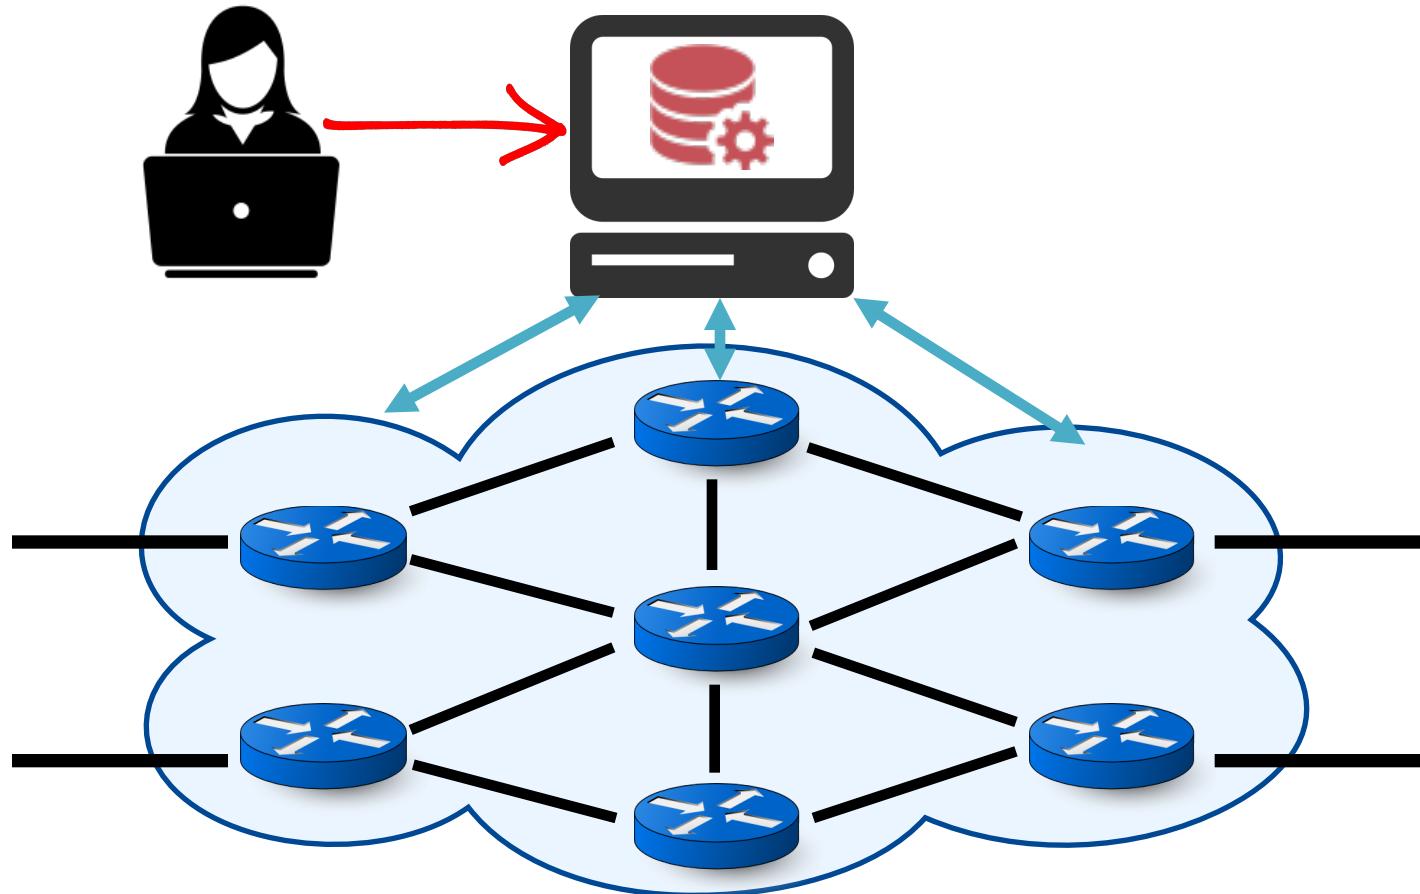
SNAP: Stateful Network-Wide Abstractions for Packet Processing

Mina Tahmasbi Arashloo¹, **Yaron Koral**¹, Michael Greenberg²,
Jennifer Rexford¹, and David Walker¹

¹ Princeton University, ²Pomona College

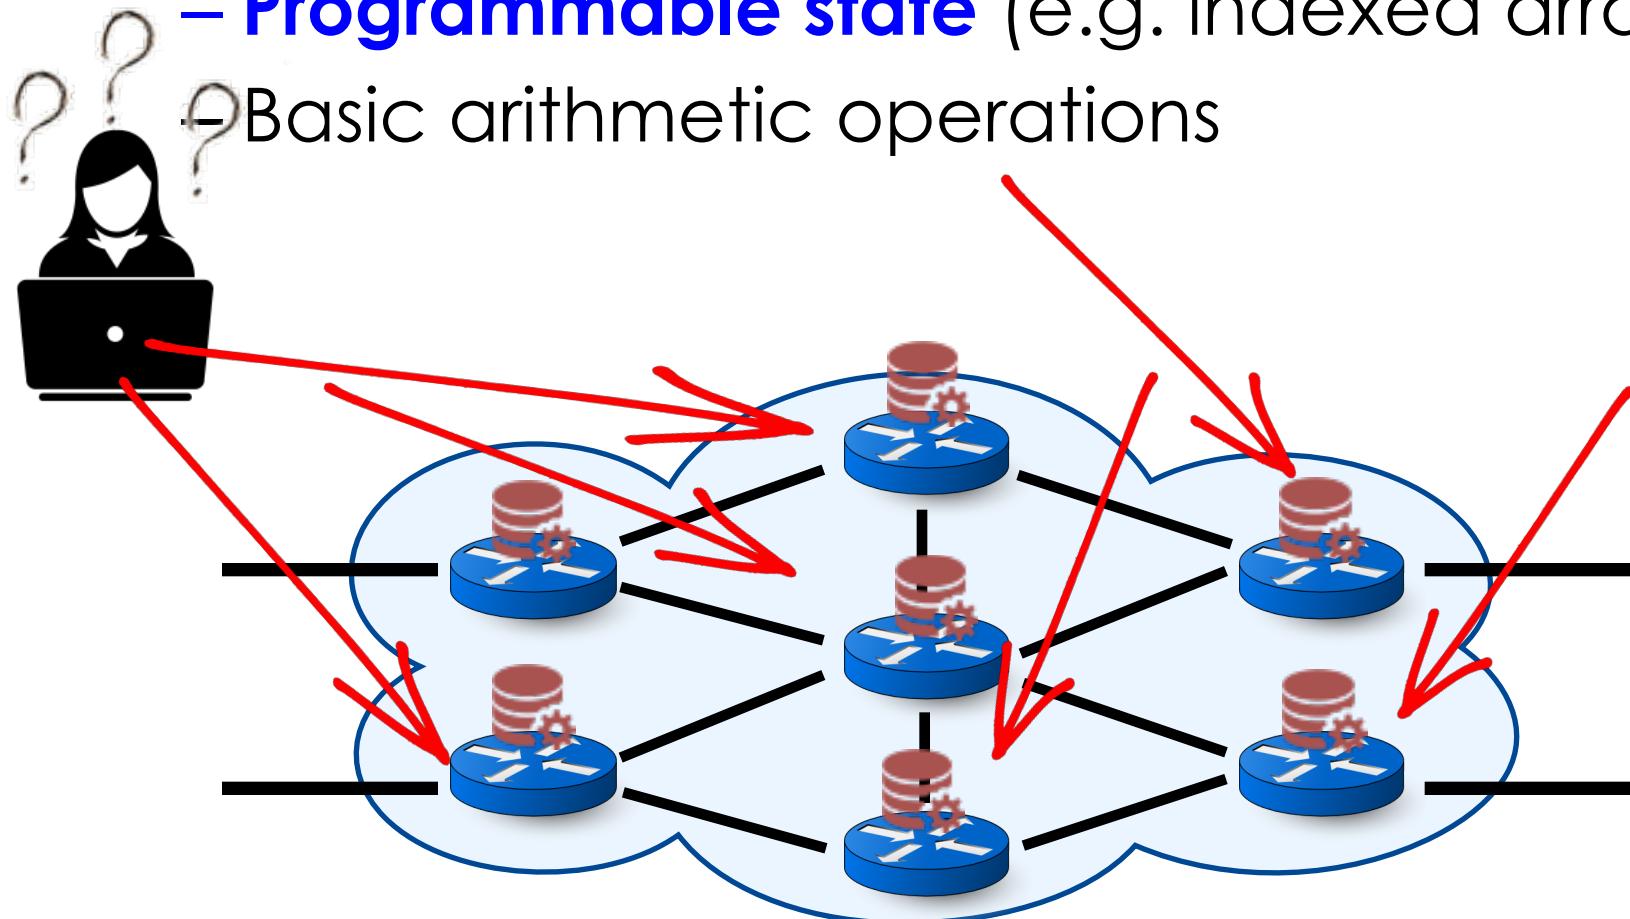
Early SDN Switch Interfaces

- Manipulate packet forwarding rules
- Read predefined set of counters

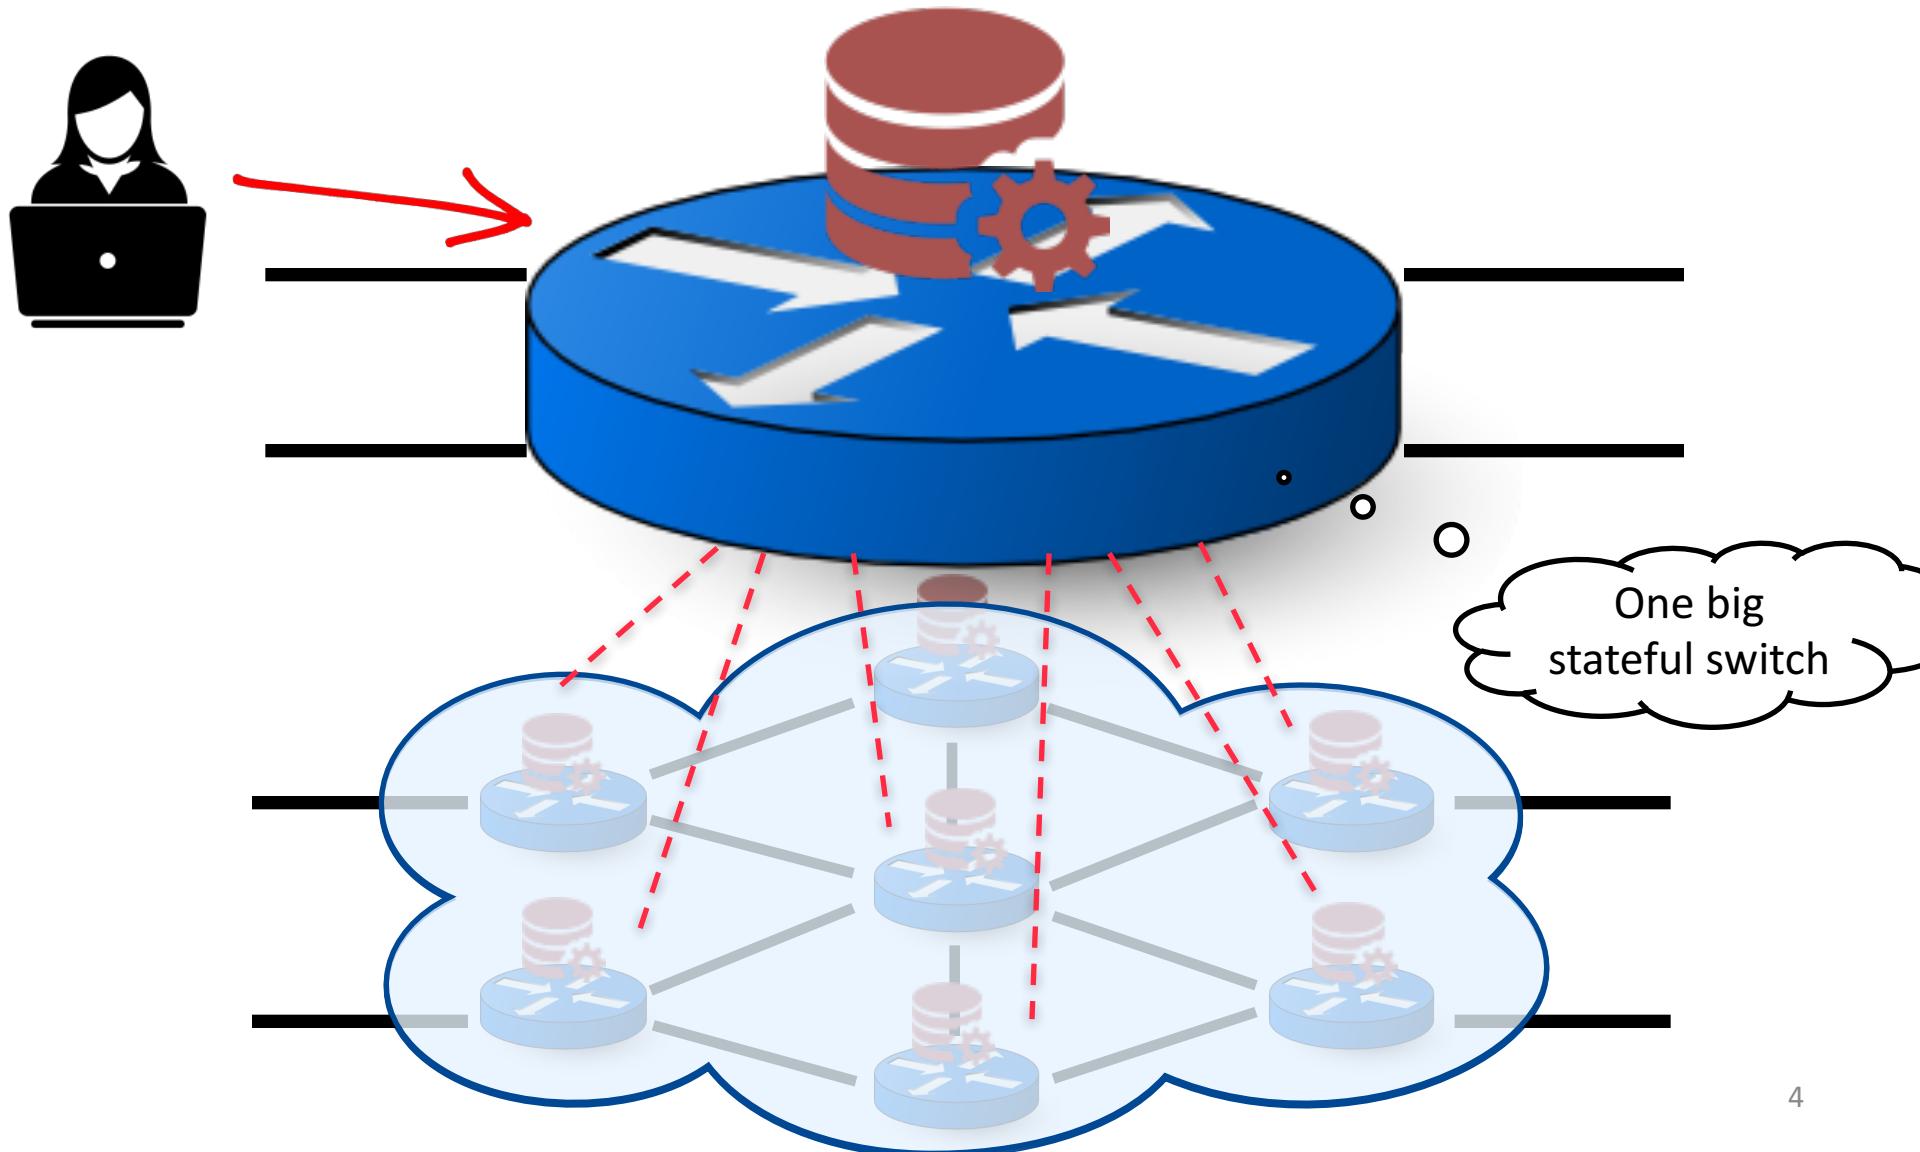


Programmable Switch Interfaces

- P4, OpenState, Open vSwitch, ...
 - **Programmable state** (e.g. indexed arrays)
 - Basic arithmetic operations



SNAP: Stateful Network Wide Programming Language



SNAP Contributions



Modular Stateful Language

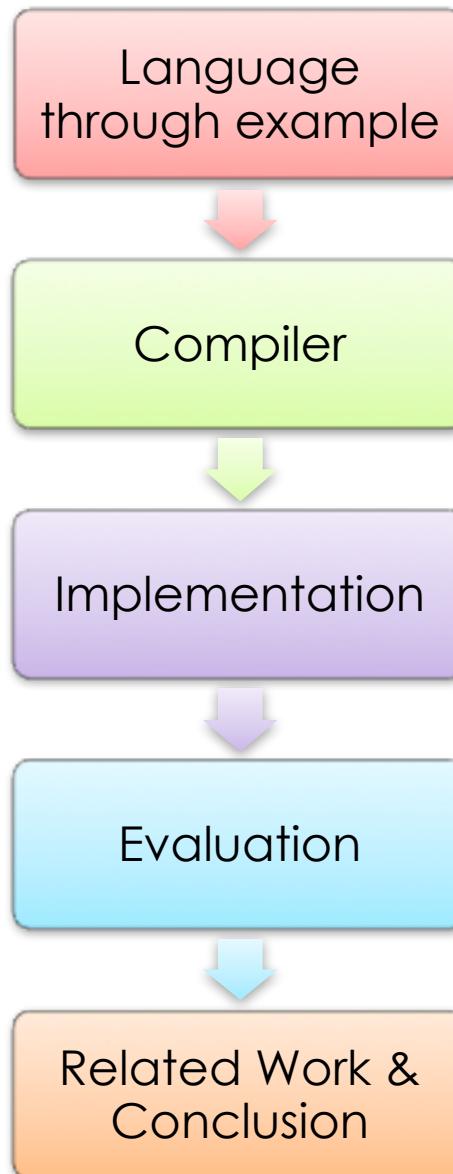


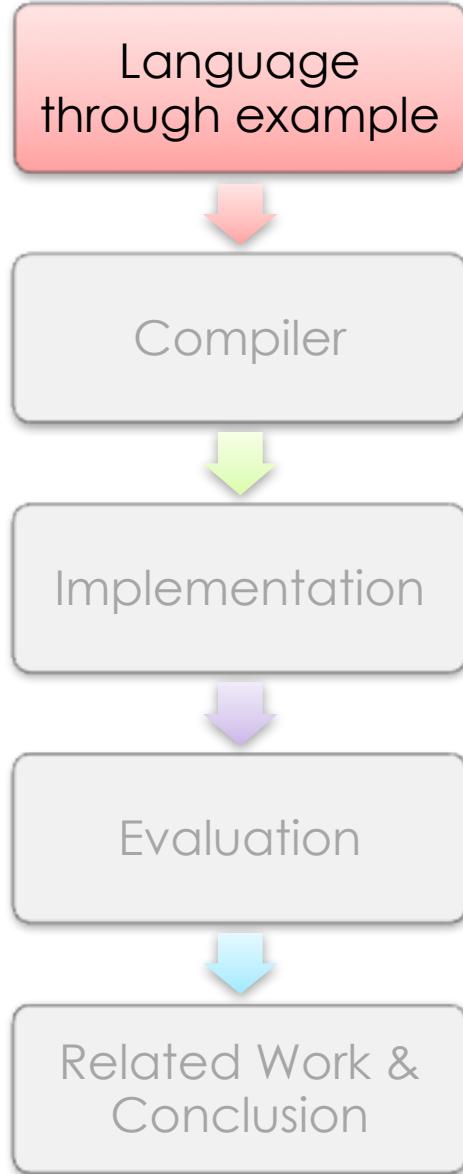
One Big Stateful Switch



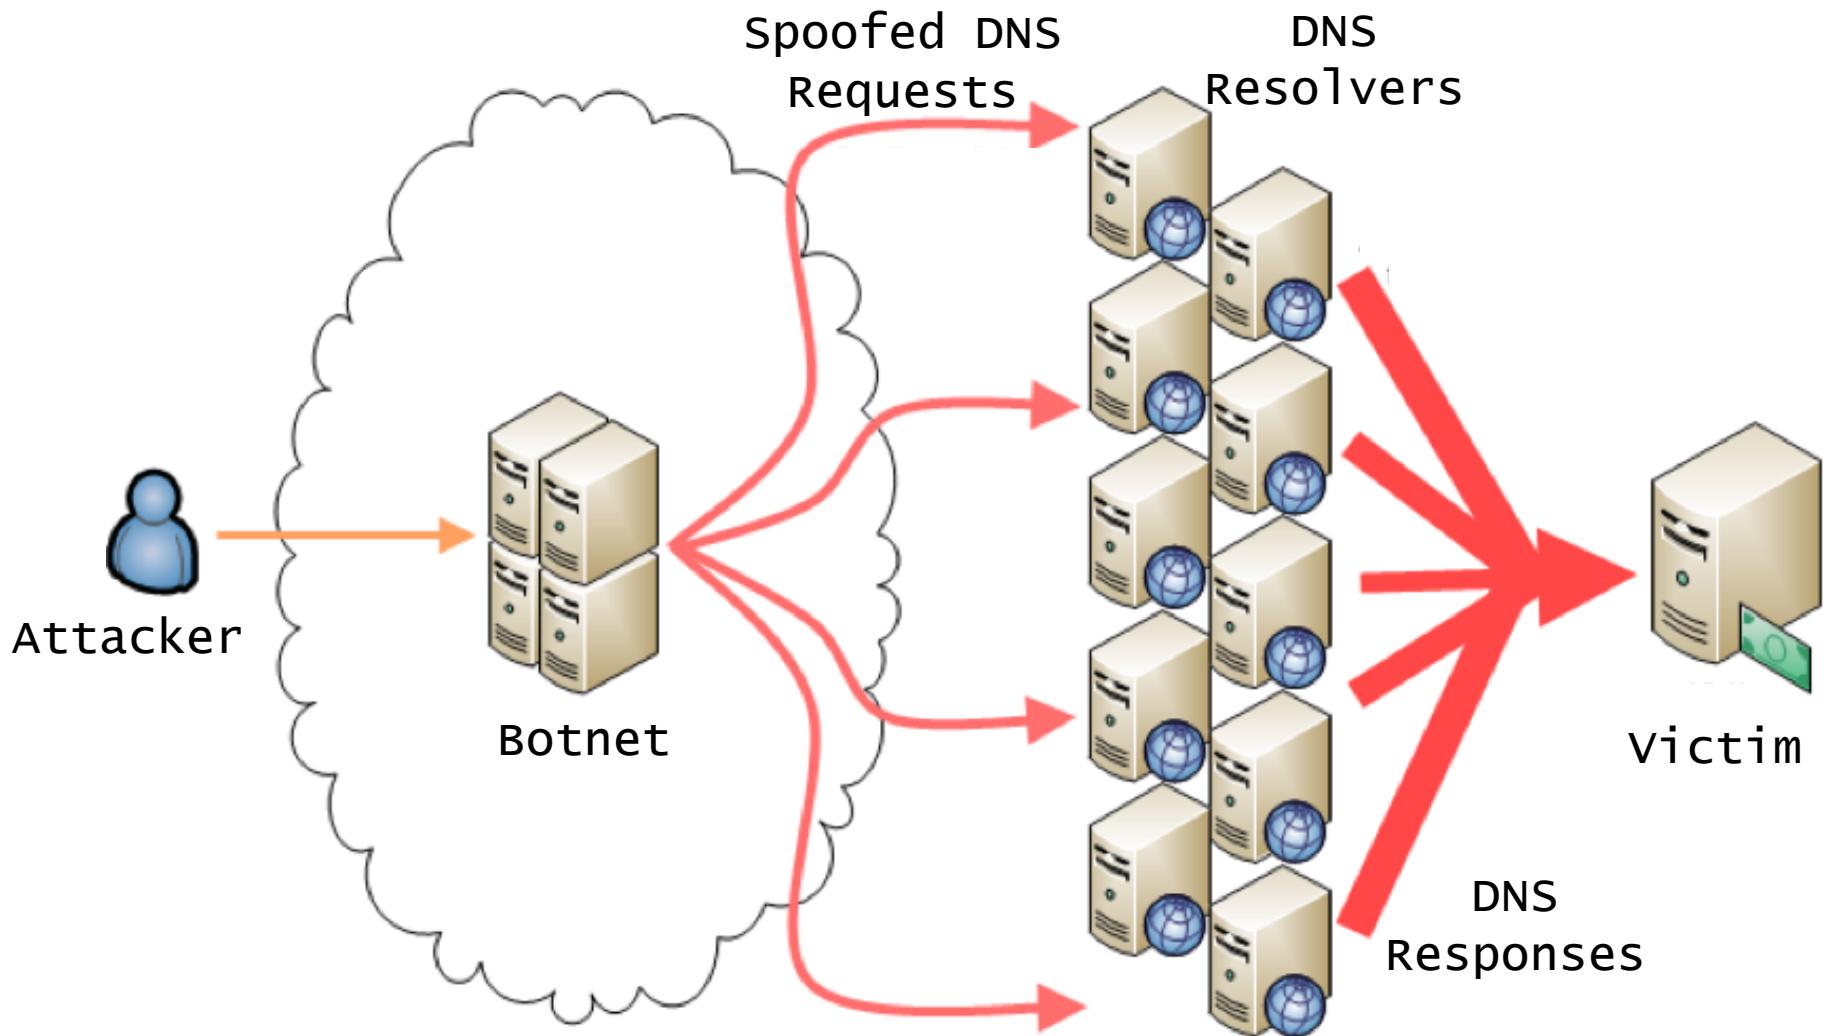
Placement + Routing

Talk Outline

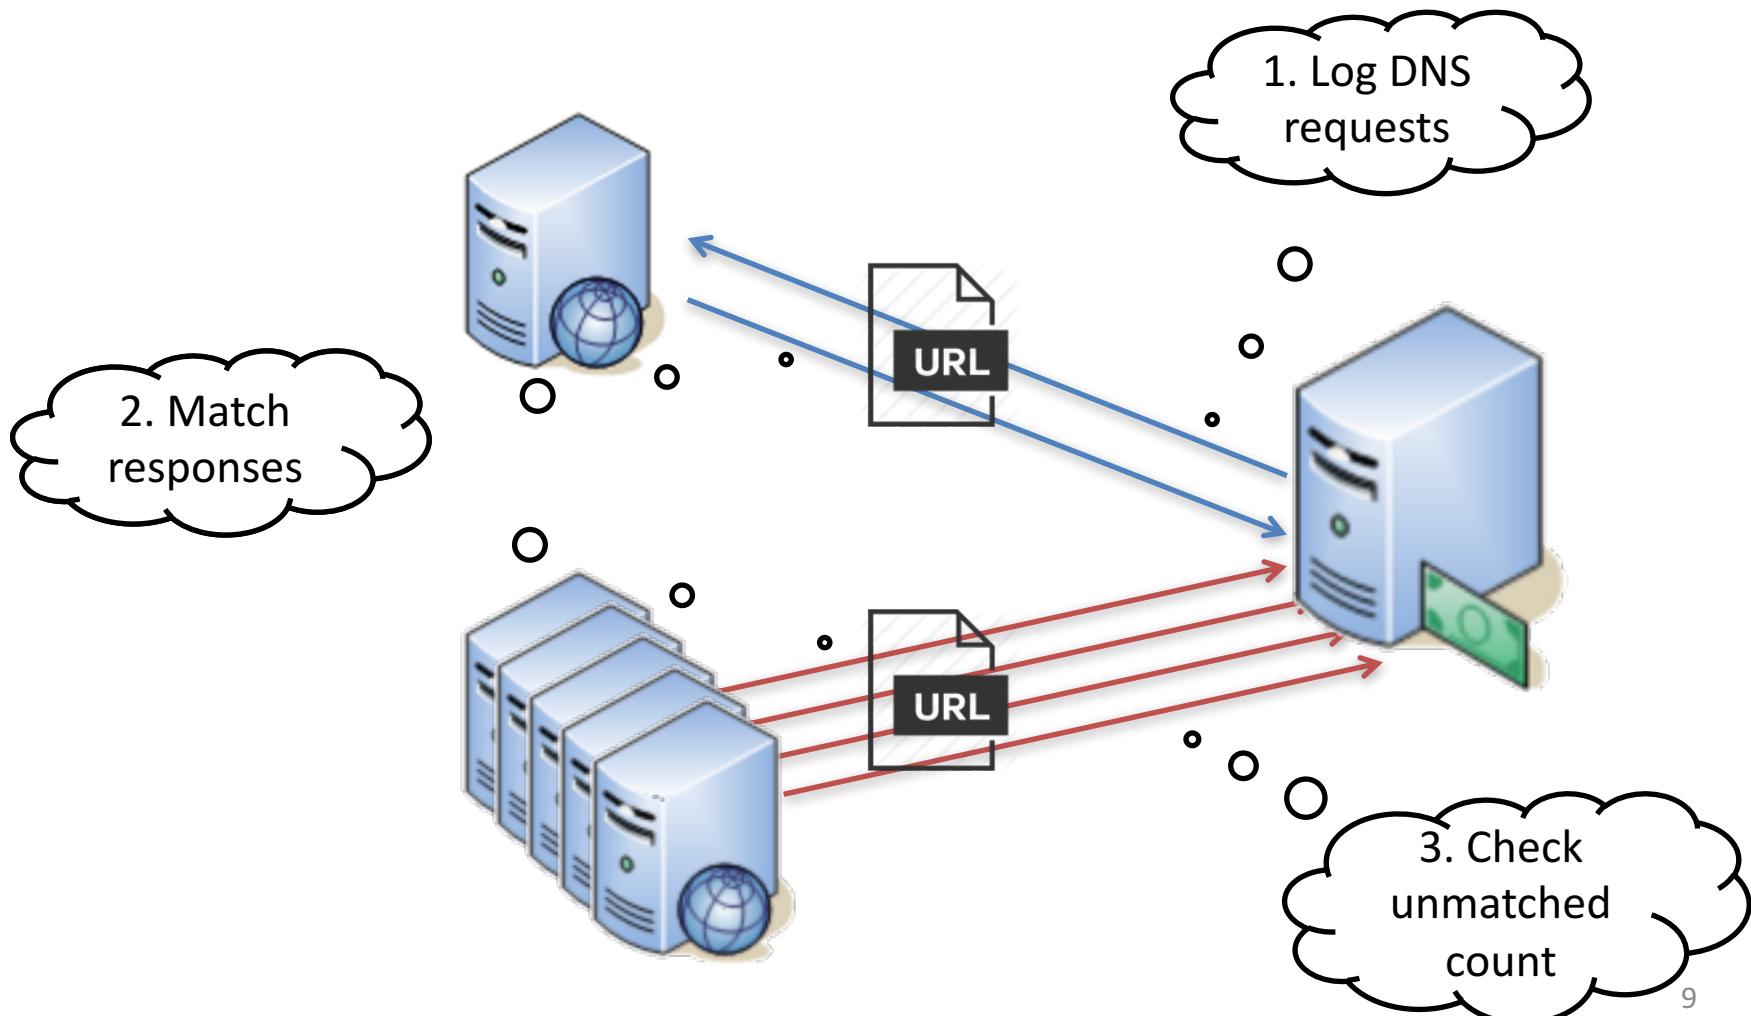




Example - DNS Reflection Attacks



Detecting DNS Reflection Attacks



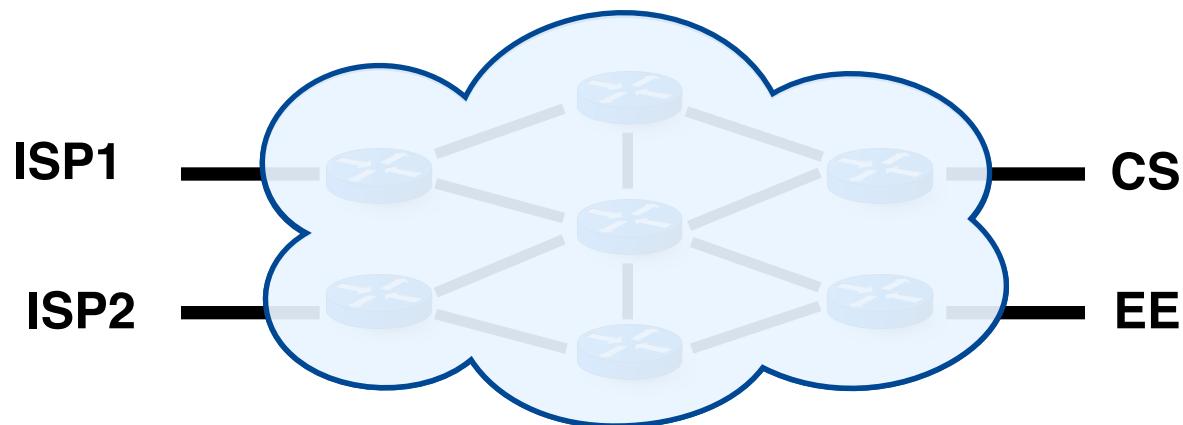
DNS Reflection Detection in SNAP

```
if srcip in CSNET & dstport = 53 then
    seen[srcip][dns.id] ← True
else if dstip in CSNET & srcport = 53 then
    if ~seen[dstip][dns.id] then
        unmatched[dstip]++;
        if unmatched[dstip] = threshold then
            susp[dstip] ← True
    else id
else id
```

- **Seen**: Keep track of DNS requests by client and DNS identifier
- **Unmatched**: Count DNS responses that don't match prior requests
- **Susp**: Suspected victims receive many unmatched responses

OBSS Forwarding in SNAP

```
if dstip = CSNET then outport ← CS  
else if dstip = EENET then outport ← EE  
else if dstip = ISP1NET then outport ← ISP1  
else if dstip = ISP2NET then outport ← ISP2  
else drop
```



Single Network Policy

DNS Reflection Detection

```
if srcip in CSNET & dstport = 53 then
    seen[srcip][dns.id] ← True
else if dstip in CSNET & srcport = 53 then
    if ~seen[dstip][dns.id] then
        unmatched[dstip]++;
        if unmatched[dstip] = threshold then
            susp[dstip] ← True
    else id
else id
```

Forwarding

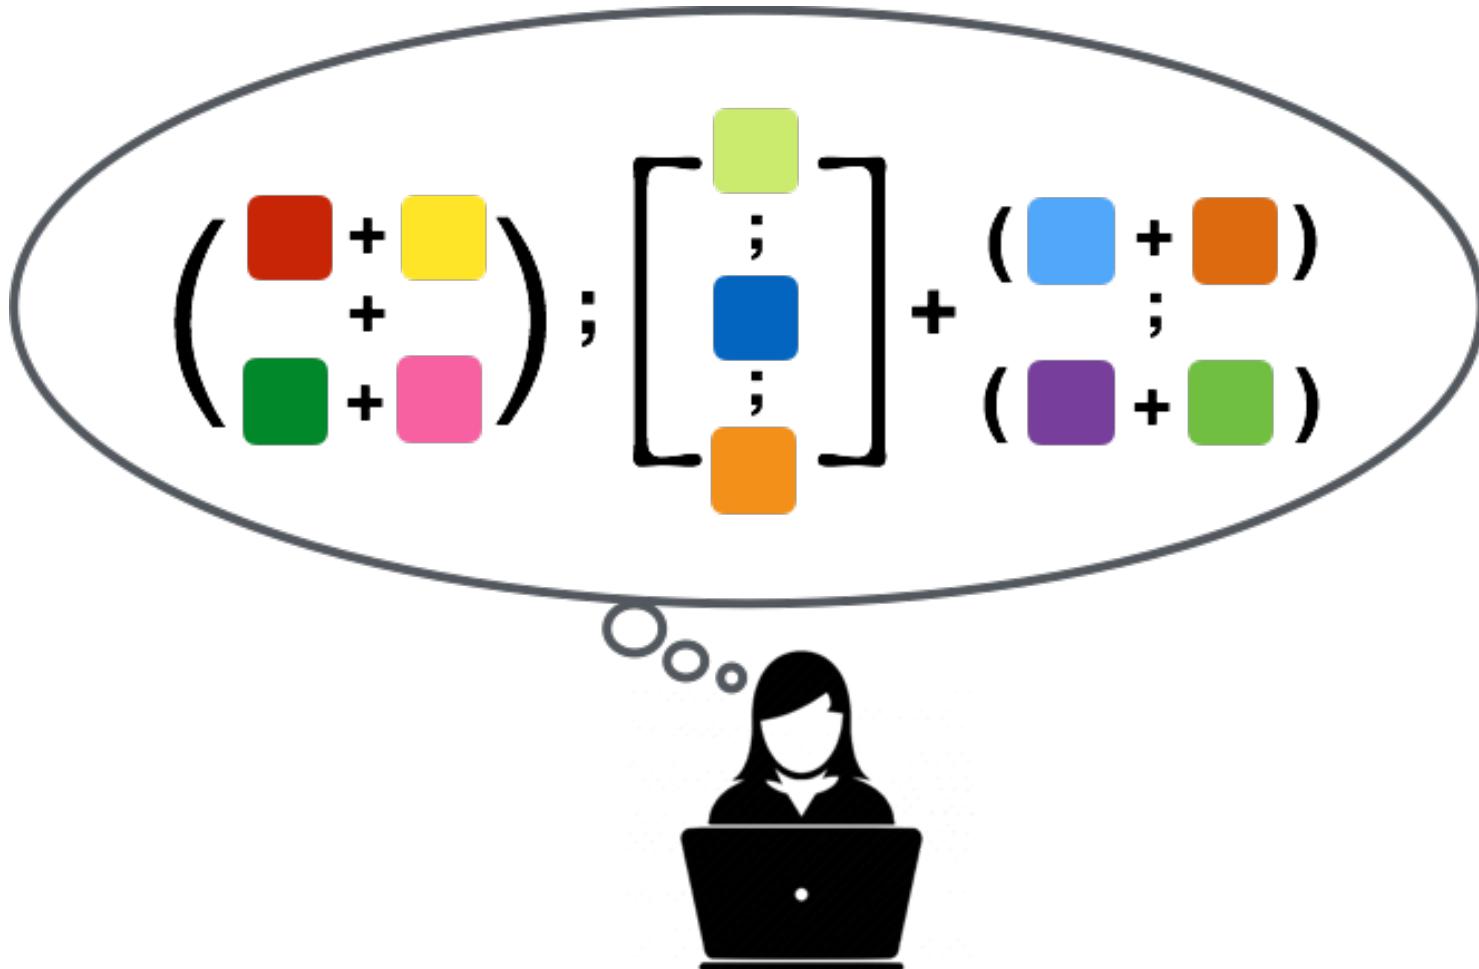
```
if dstip = CSNET then outport ← CS
else if dstip = EENET then outport ← EE
else if dstip = ISP1NET then outport ← ISP1
else if dstip = ISP2NET then outport ← ISP2
else drop
```

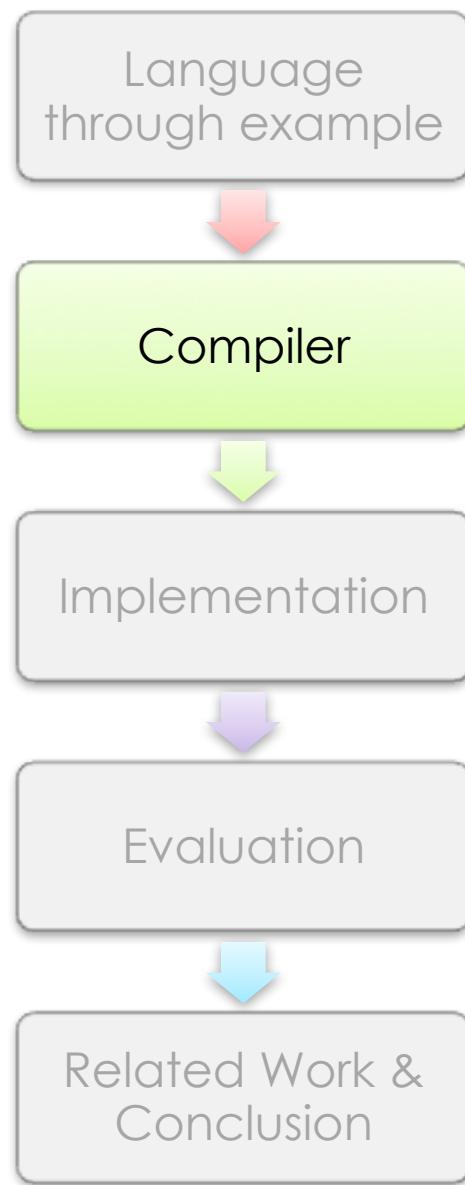


SNAP Applications

Source	Application
Chimera (USENIX Security'12)	Number of domains sharing the same IP address Number of distinct IP addresses under the same domain DNS TTL change tracking DNS tunnel detection Sidejack detection Phishing/spam detection
FAST (HotSDN'14)	Stateful firewall FTP monitoring Heavy-hitter detection Super-spreader detection Sampling based on flow size Selective packet dropping (MPEG frames) Connection affinity
Bohatei (USENIX Security'15)	SYN flood detection DNS reflection (and amplification) detection UDP flood mitigation Elephant flows detection
Others	Bump-on-the-wire TCP state machine Snort flowbits

Single Network Policy



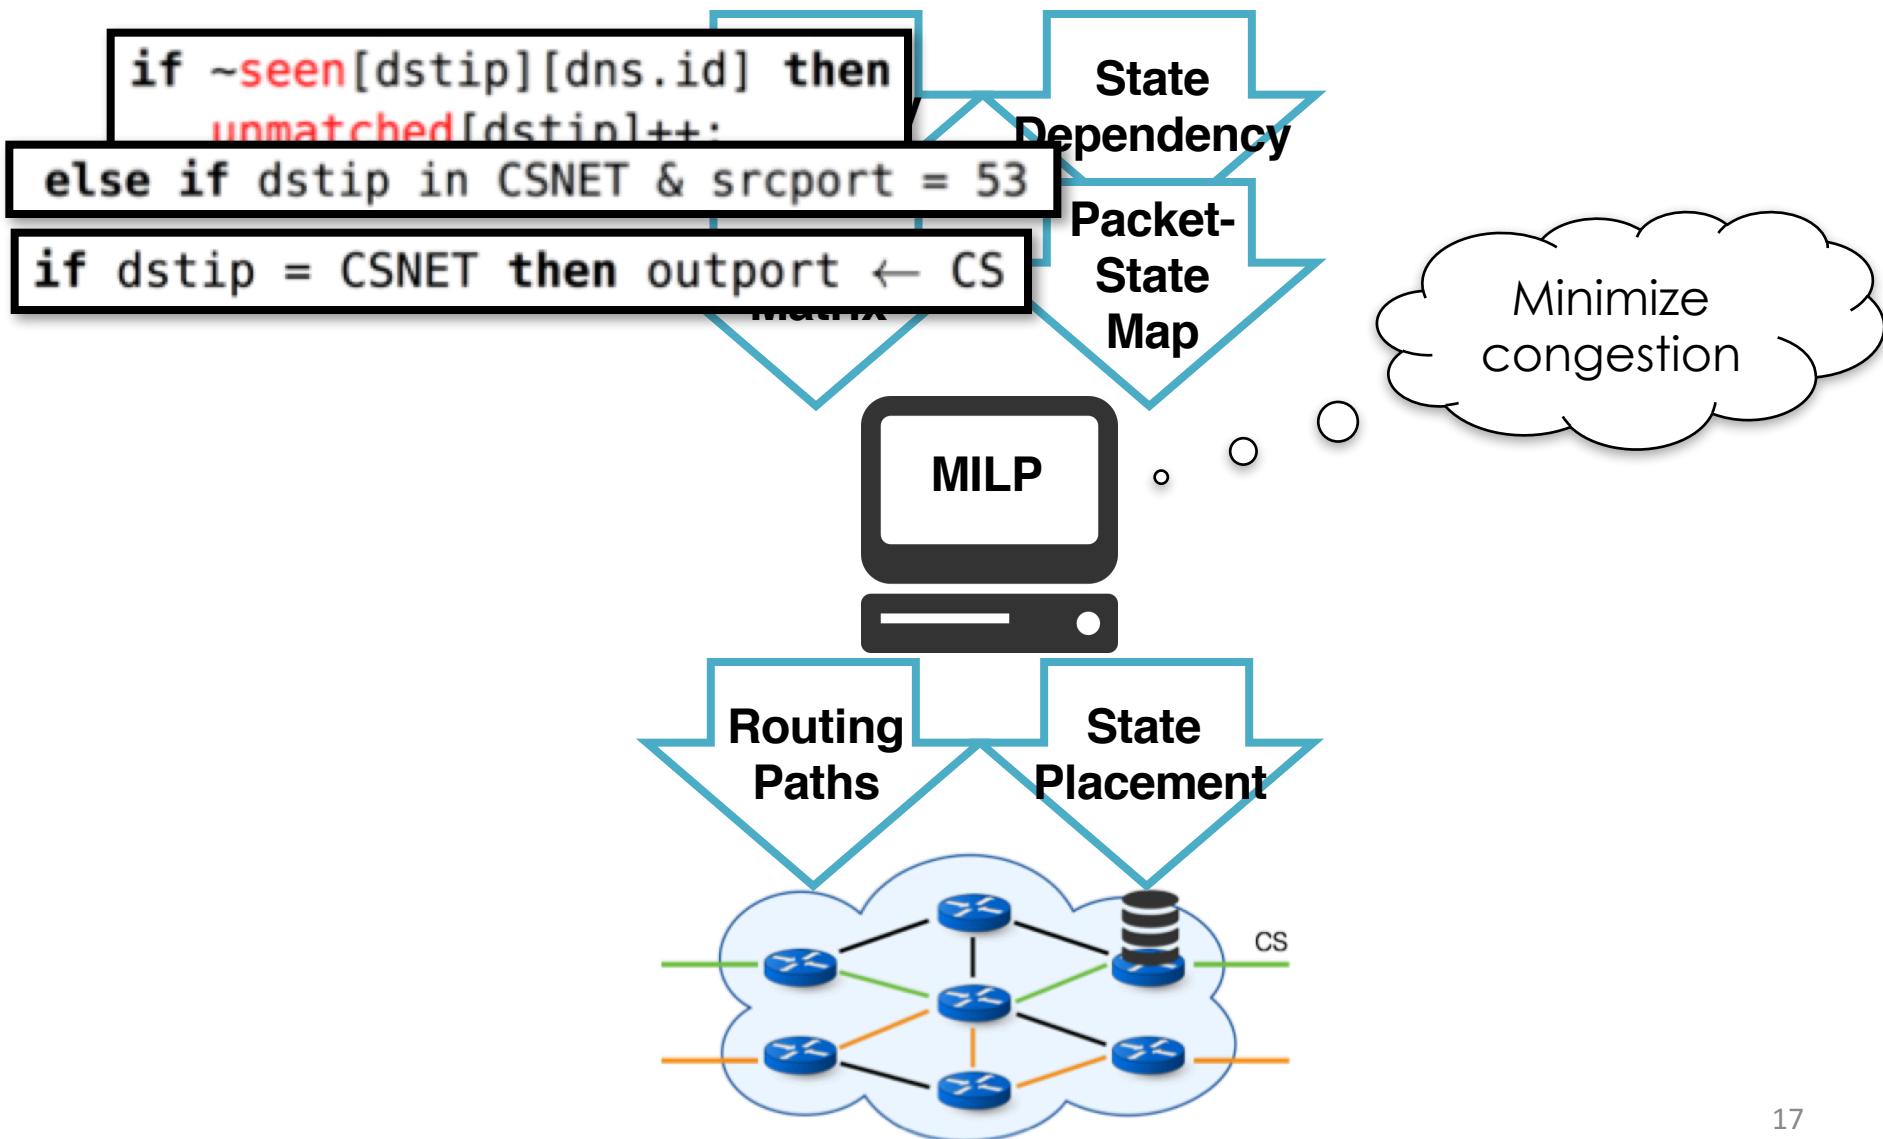


SNAP Compiler

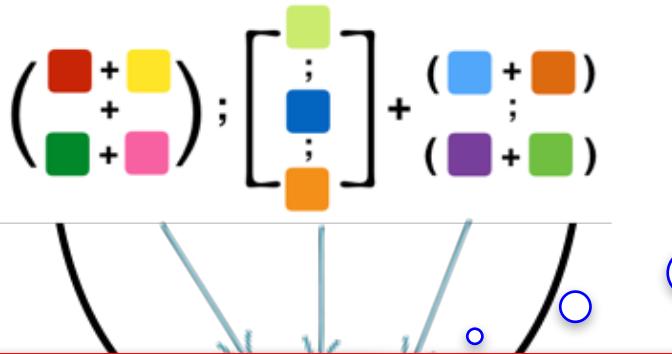
Where to place state
variables

How to forward packets
through them

Routing + Placement Jointly

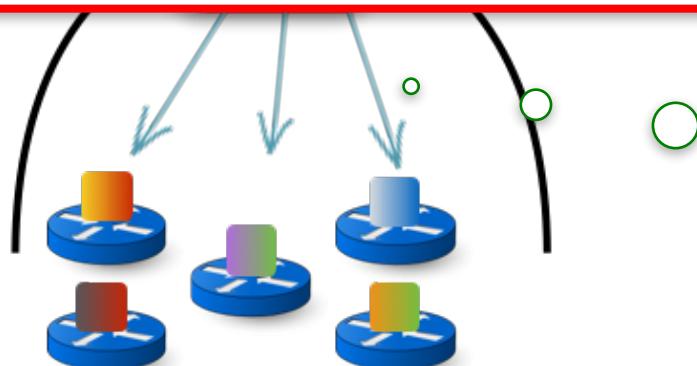


Intermediate Representation (IR)



Maintain all
programs in a single
data structure

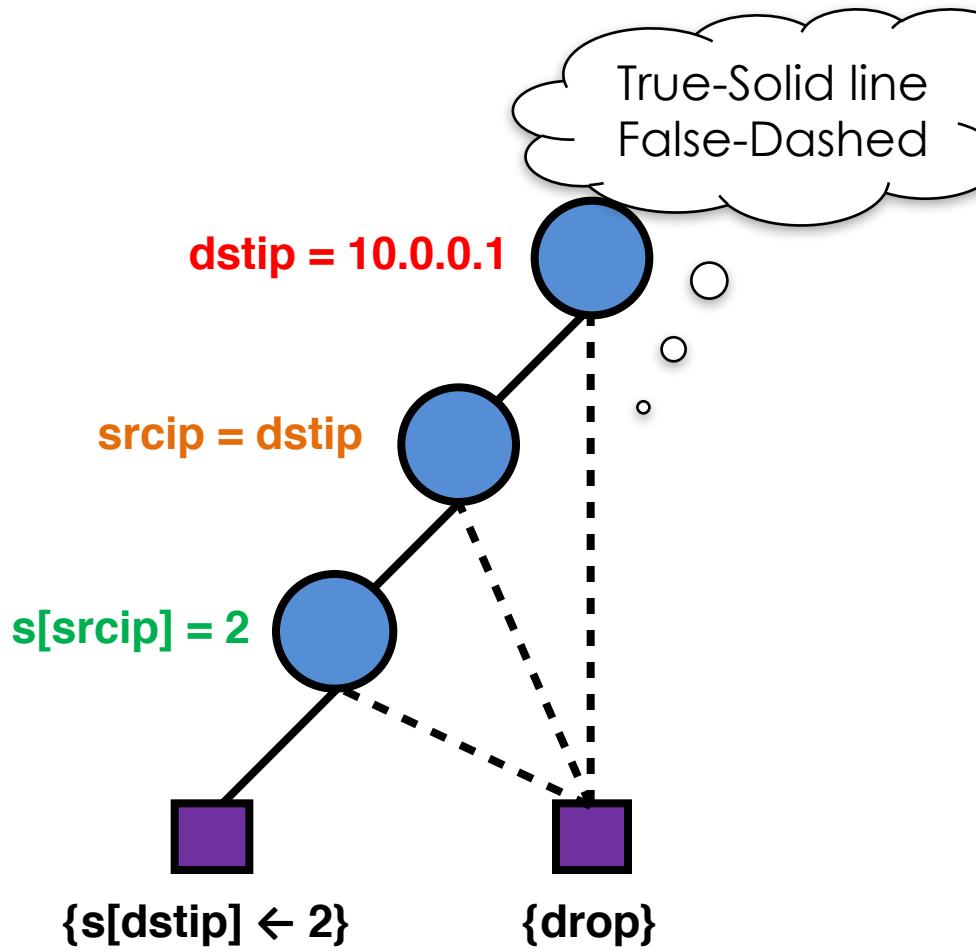
**Composable and easily
partitioned IR**



Distribute the
program to
switches

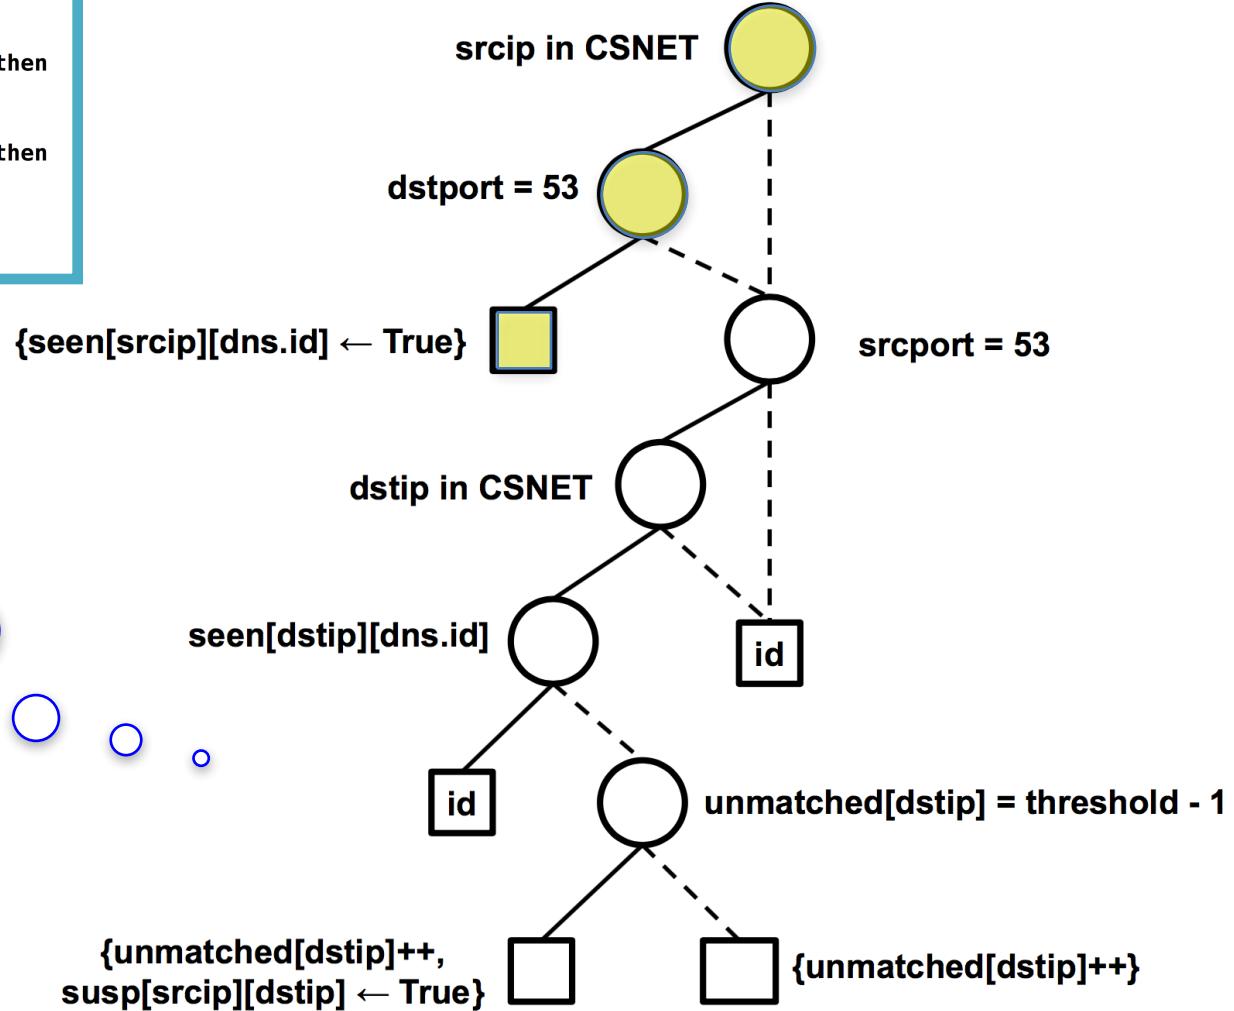
xFDDs: Extended Forwarding Decision Diagrams

- Intermediate node:
test on header
fields and state
- Leaf: set of action
sequences
- Three kinds of tests
 - $\text{field} = \text{value}$
 - $\text{field}_1 = \text{field}_2$
 - $\text{state_var}[idx] = \text{val}$



xFDD for DNS Reflection Detection

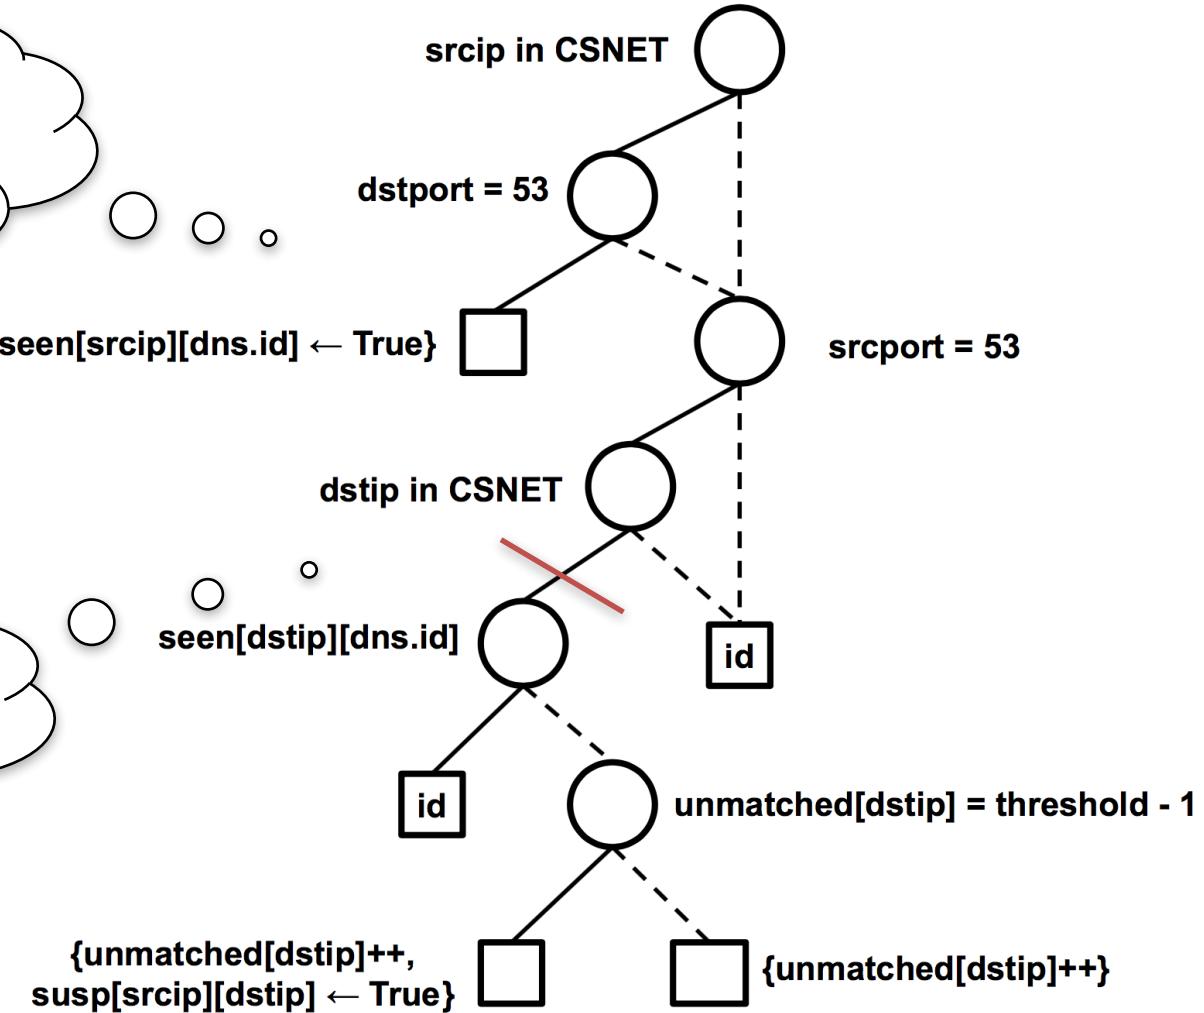
```
if srcip in CSNET & dstport = 53 then
    seen[srcip][dns.id] ← True
else if dstip in CSNET & srcport = 53 then
    if ~seen[dstip][dns.id] then
        unmatched[dstip]++;
        if unmatched[dstip] = threshold then
            susp[dstip] ← True
    else id
else id
```



xFDD for DNS Reflection Detection

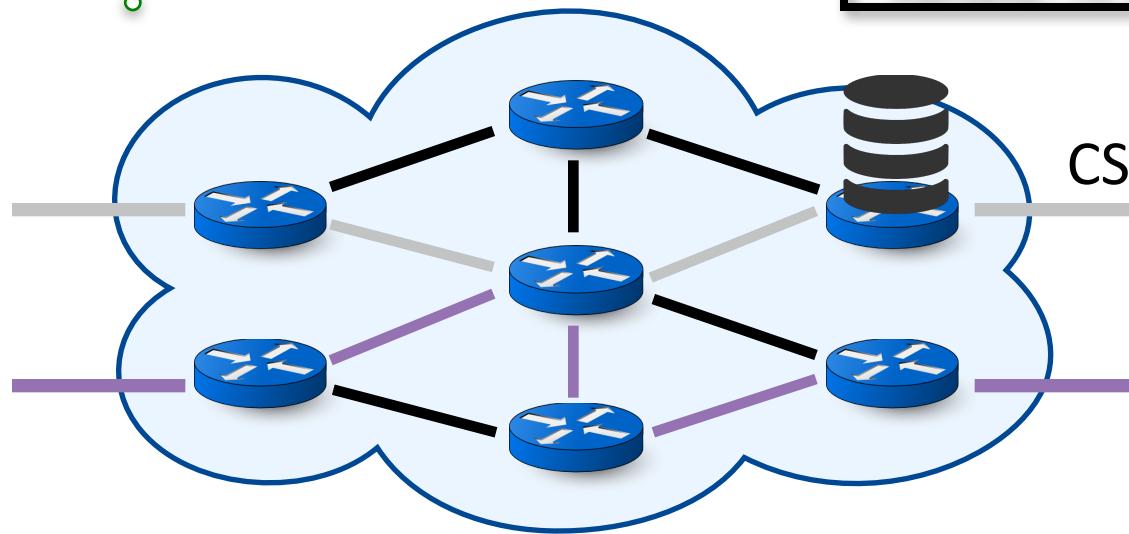
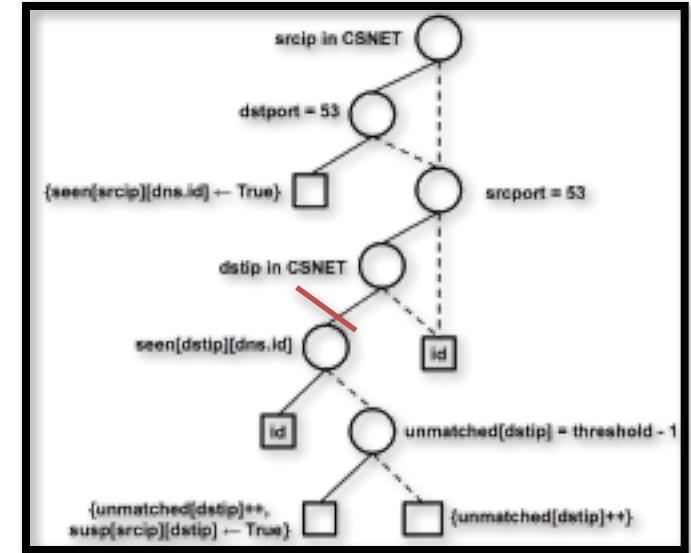
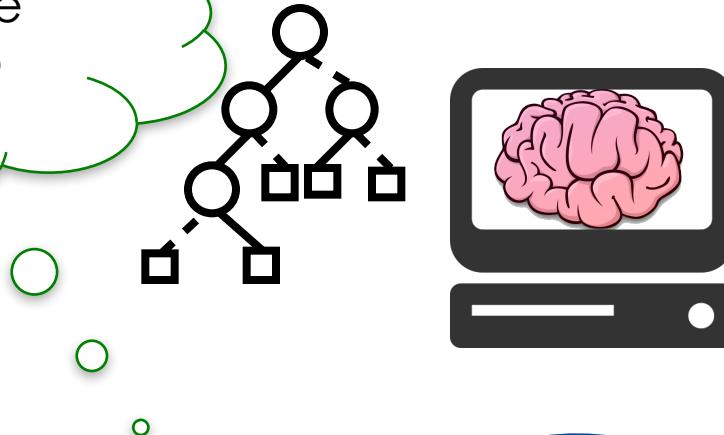
Fixes the order in which programs access state.

We could distribute the programs by placing cuts



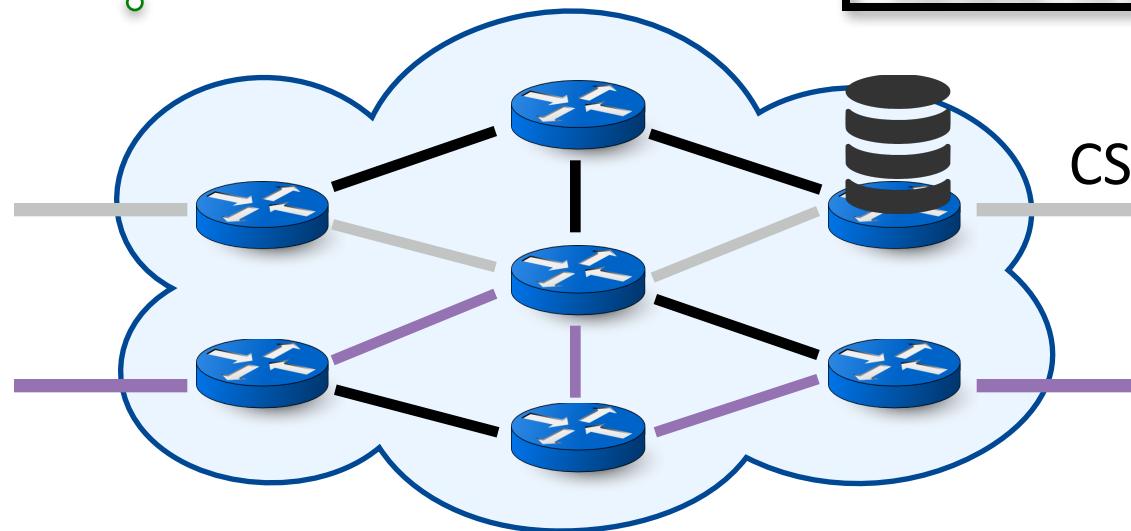
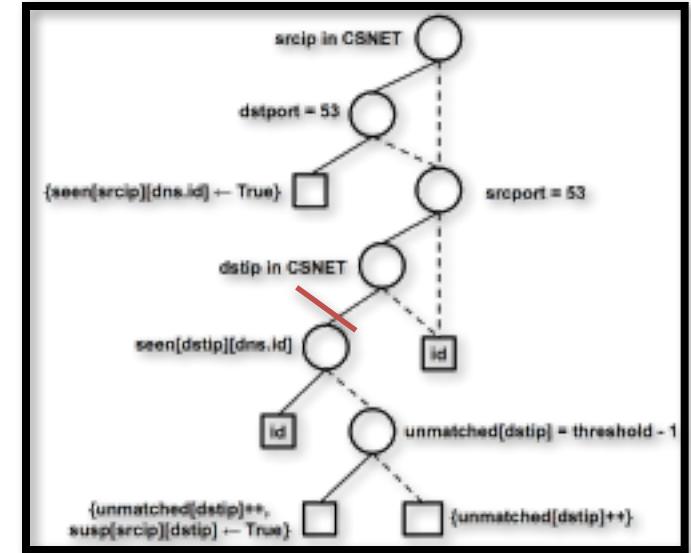
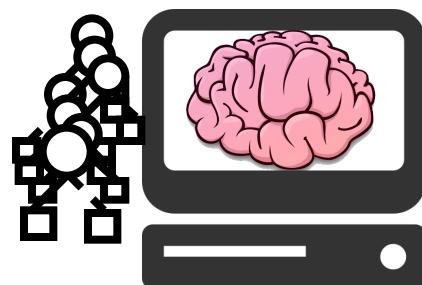
Partitioning to Sub-Programs

Distribute the program to switches

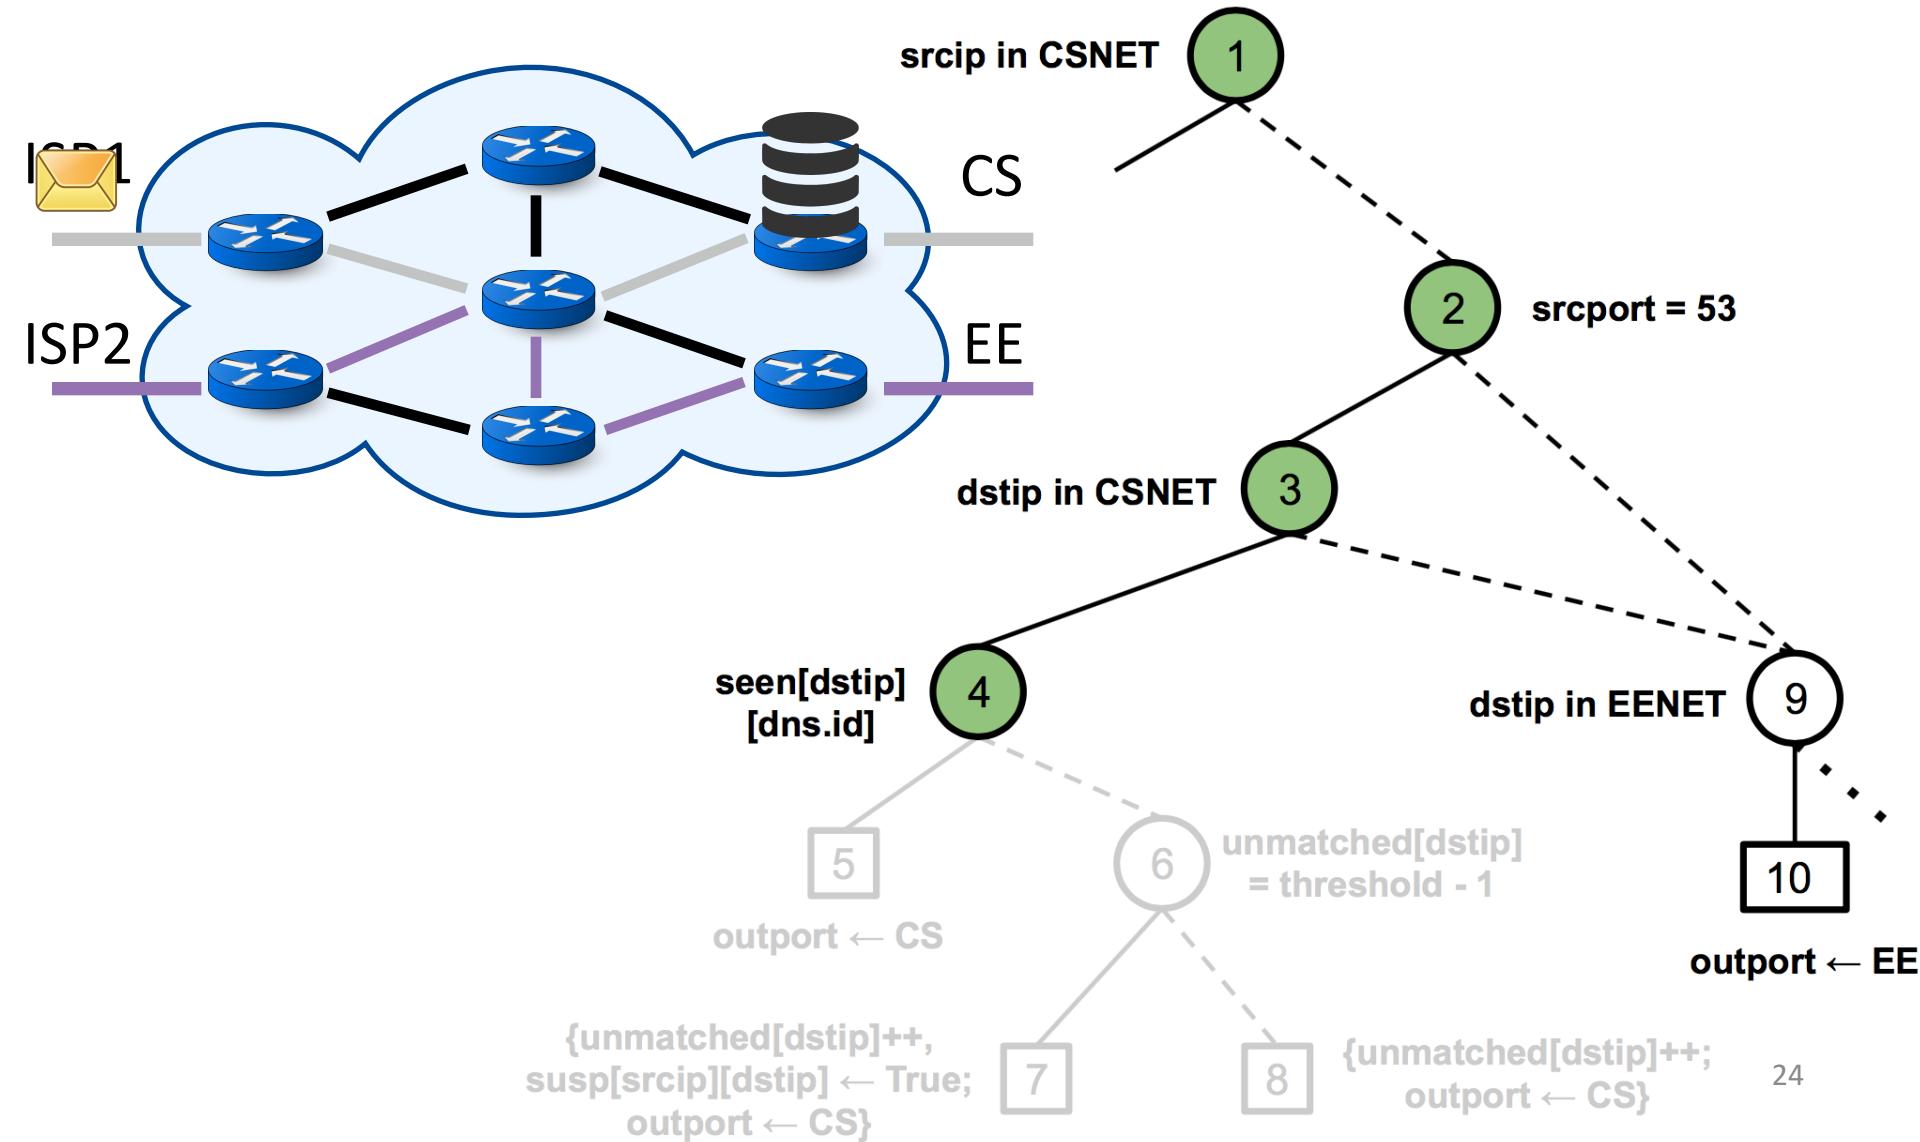


Partitioning to Sub-Programs

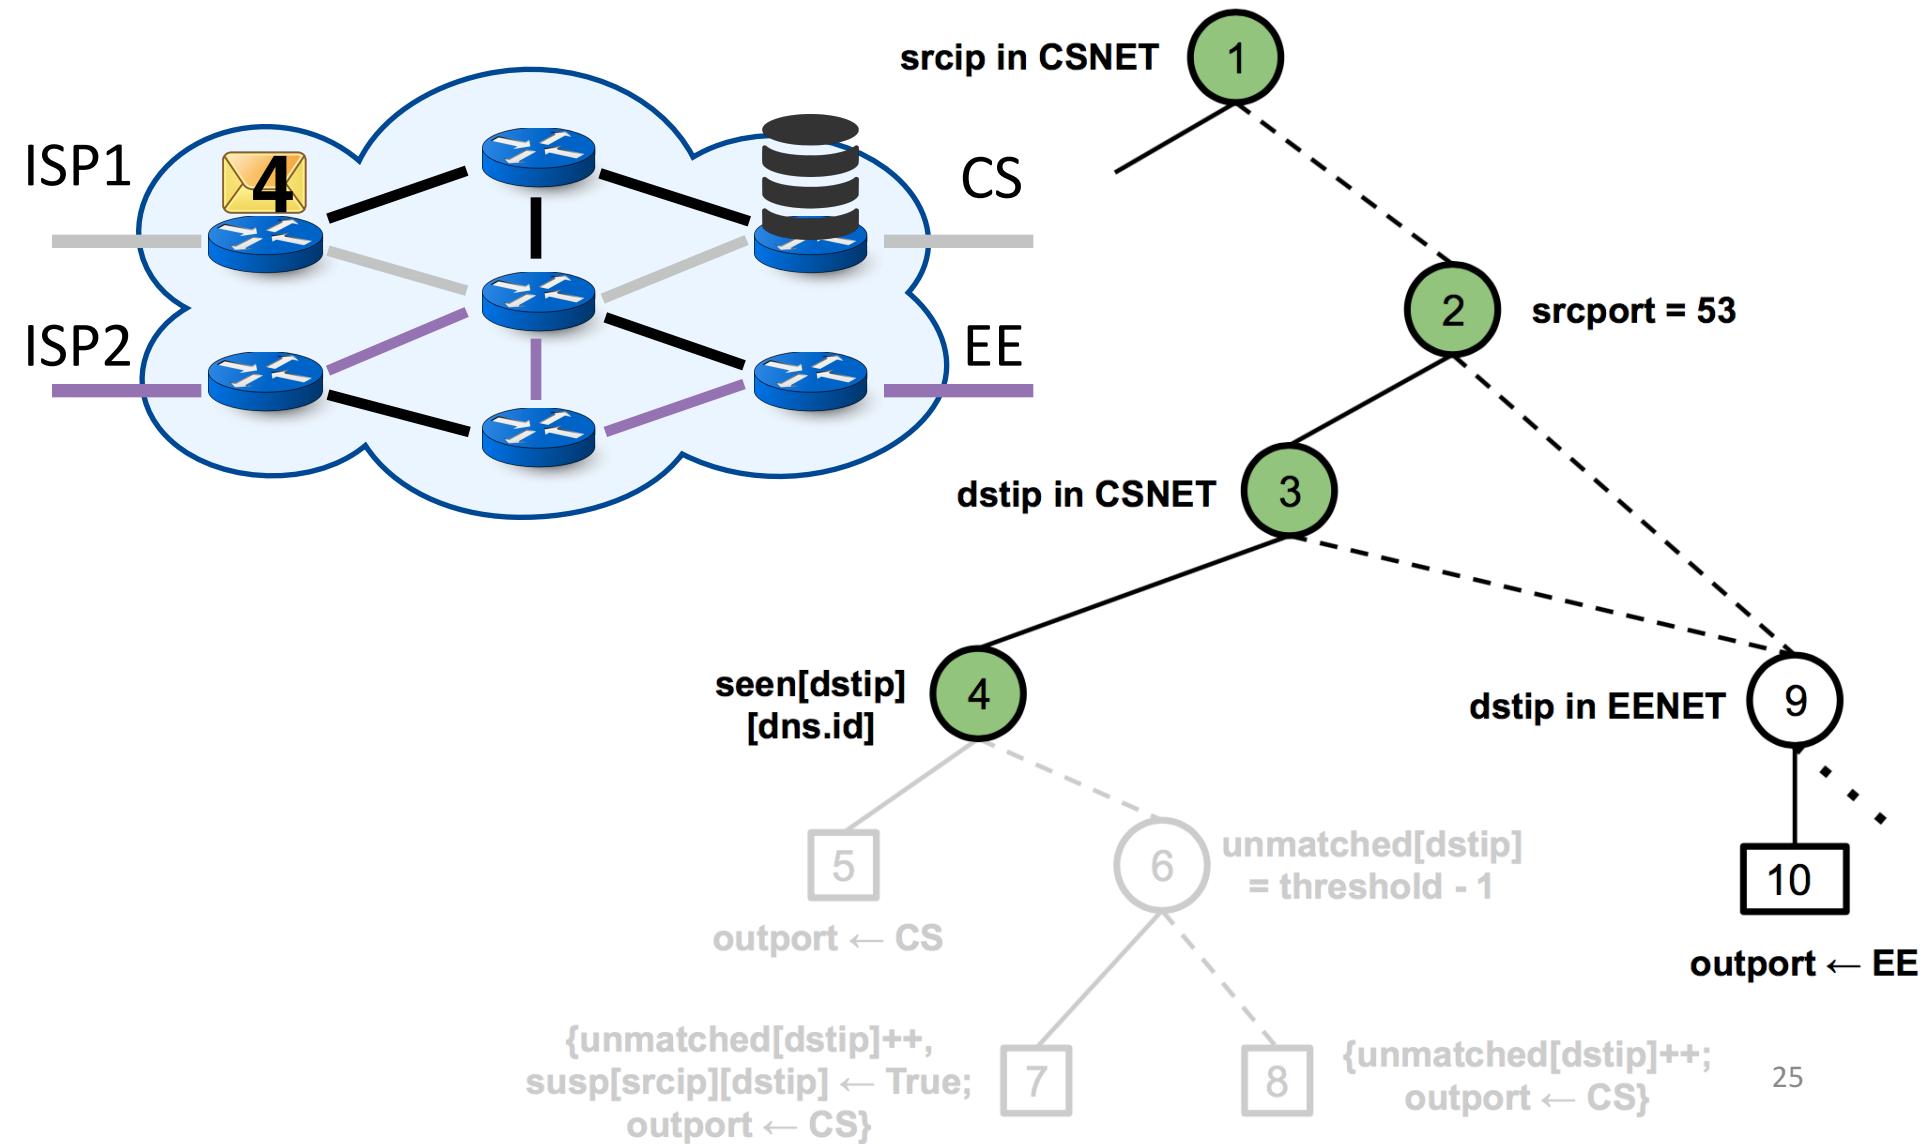
Distribute the program to switches



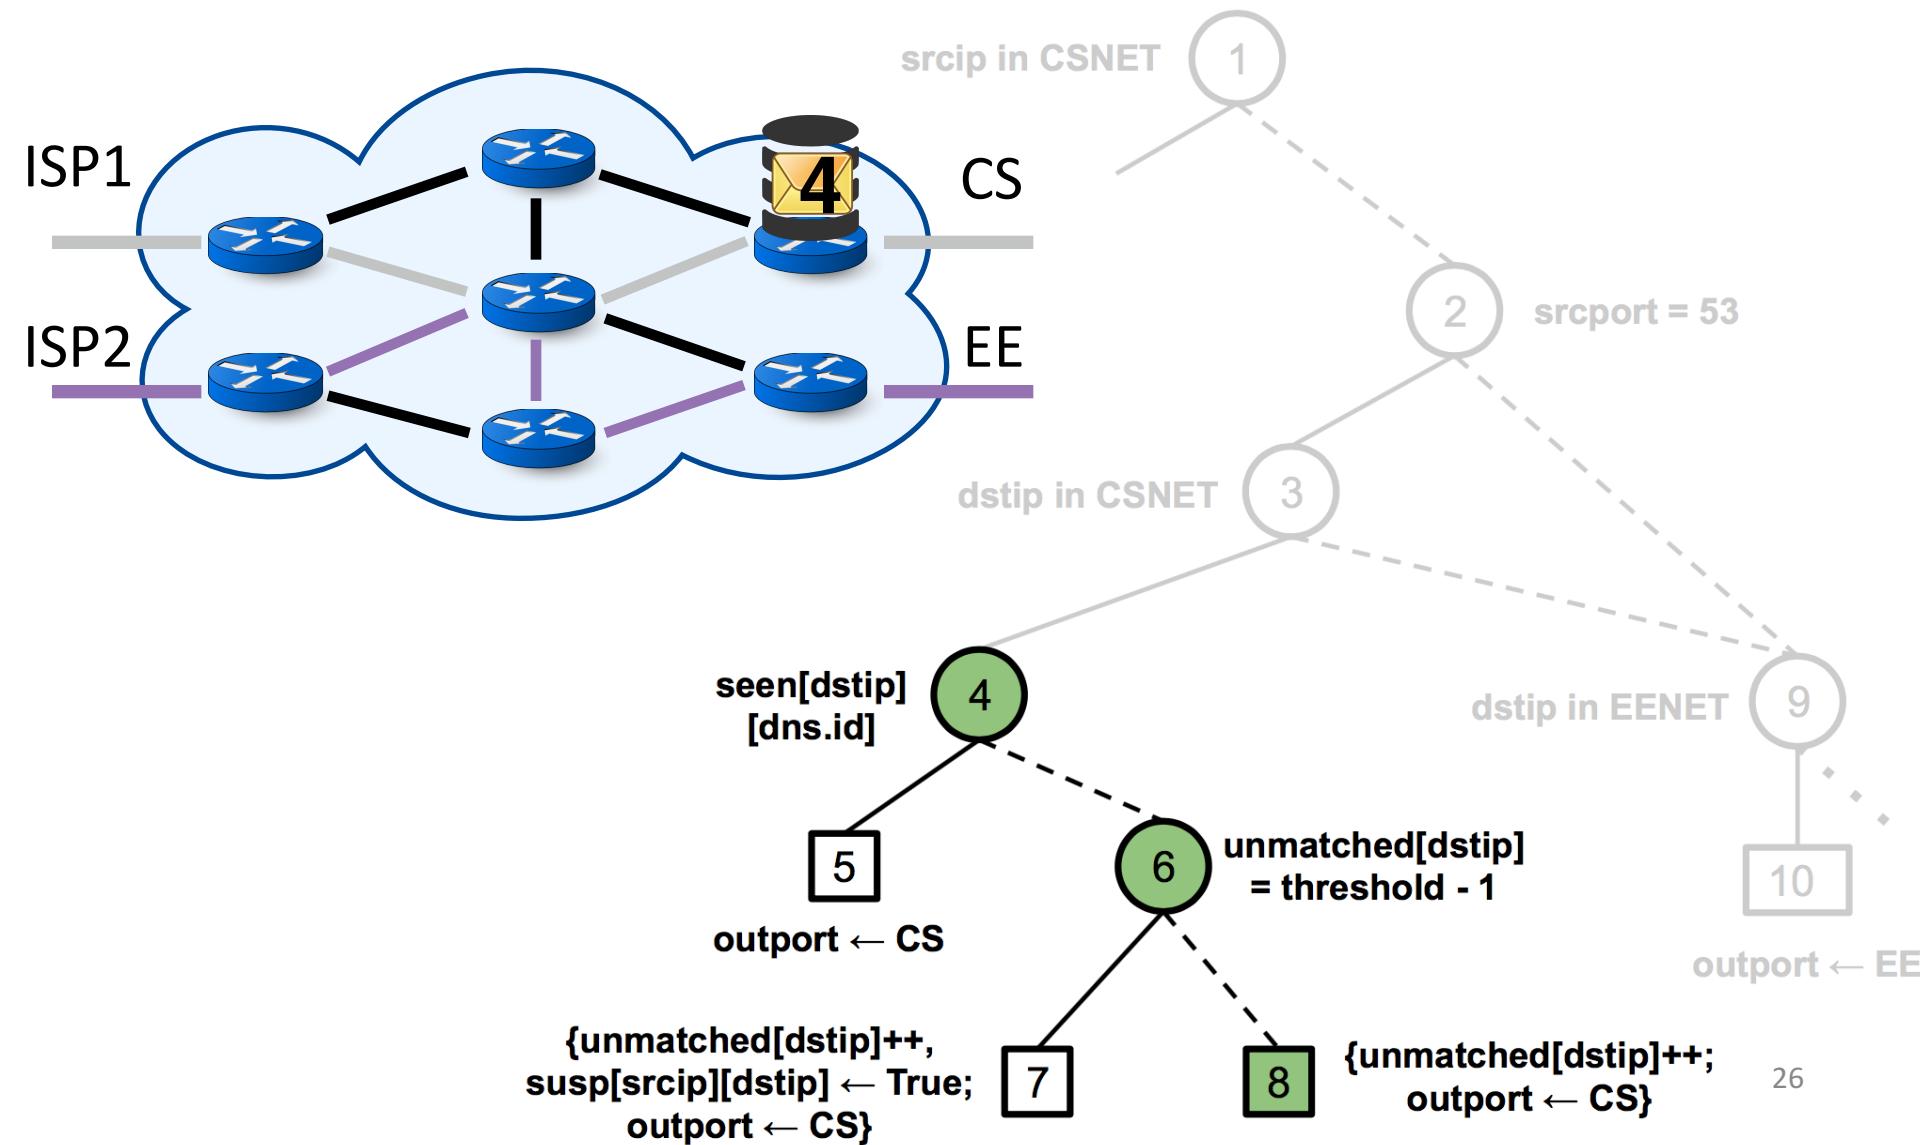
Putting It All Together



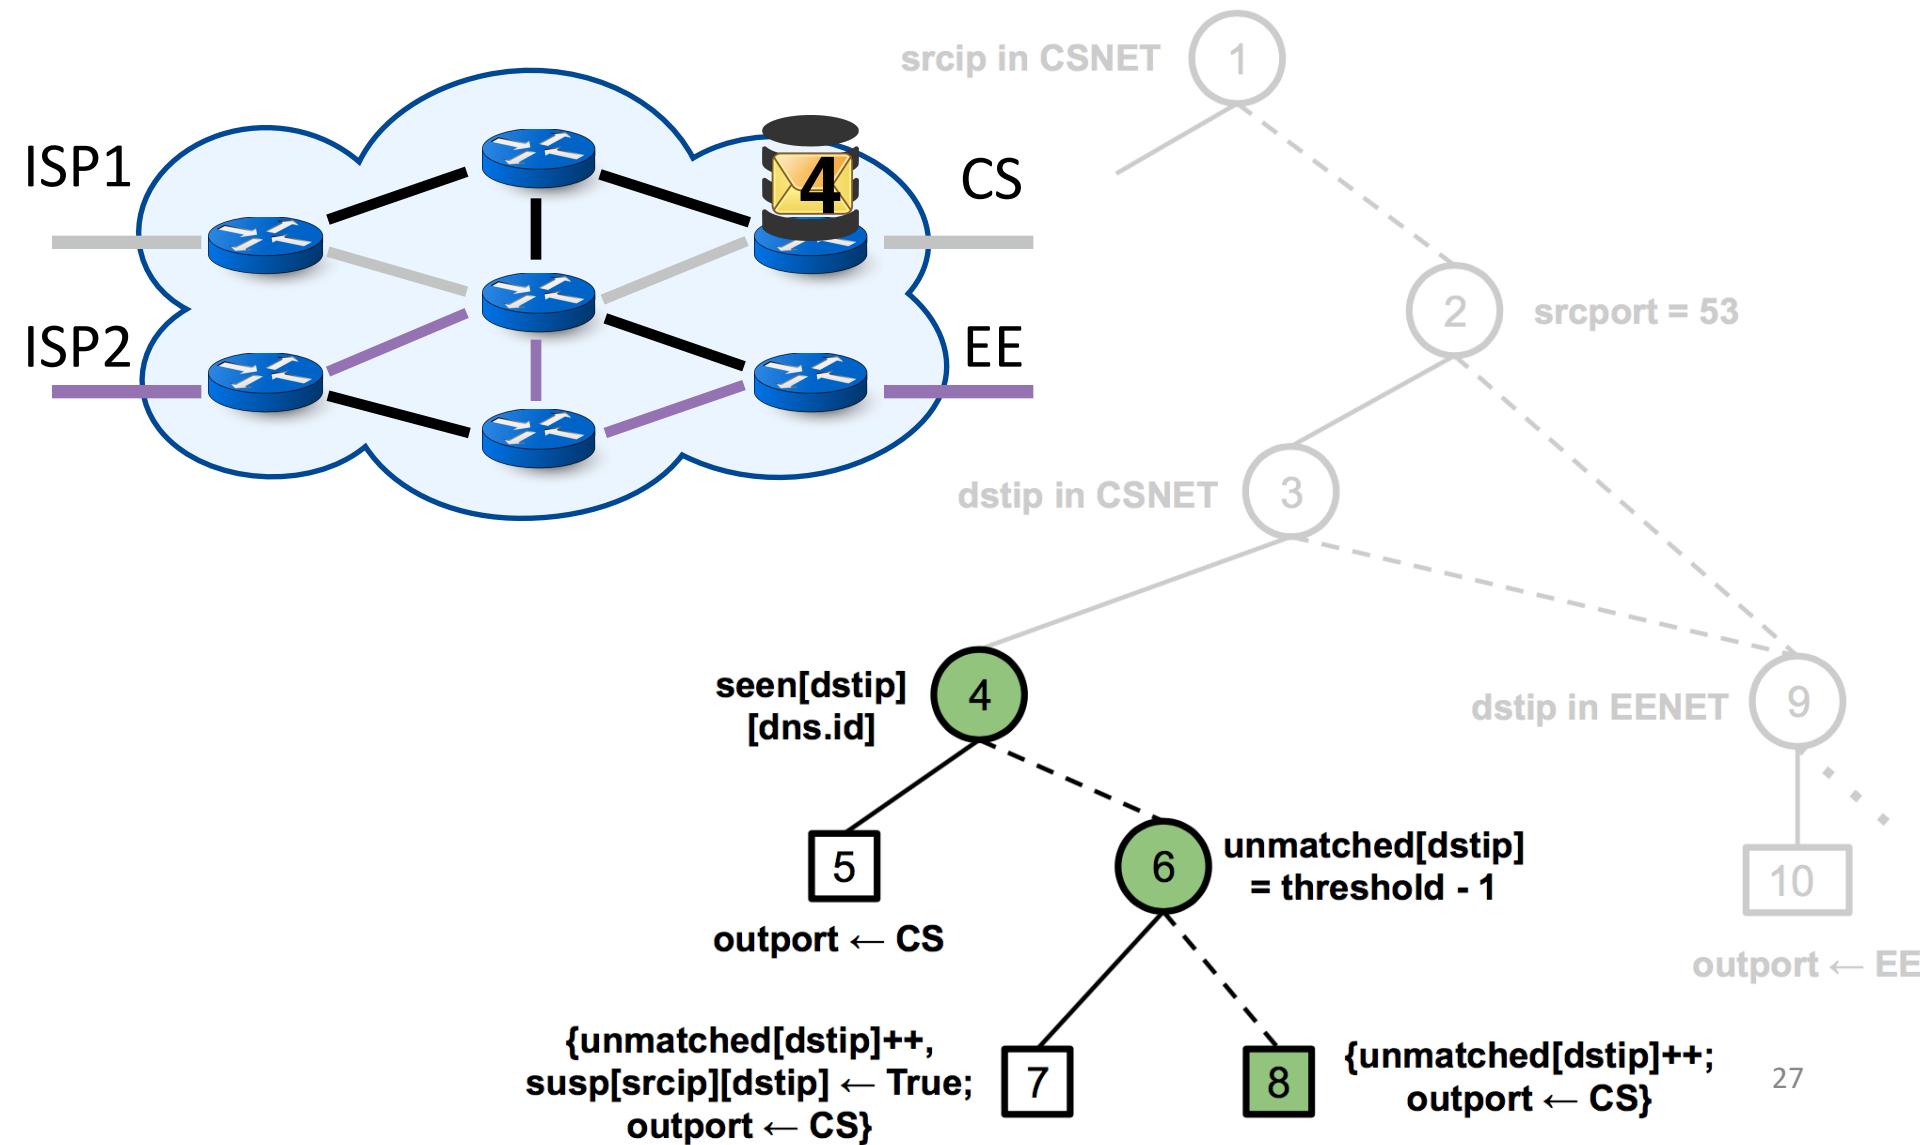
Putting It All Together

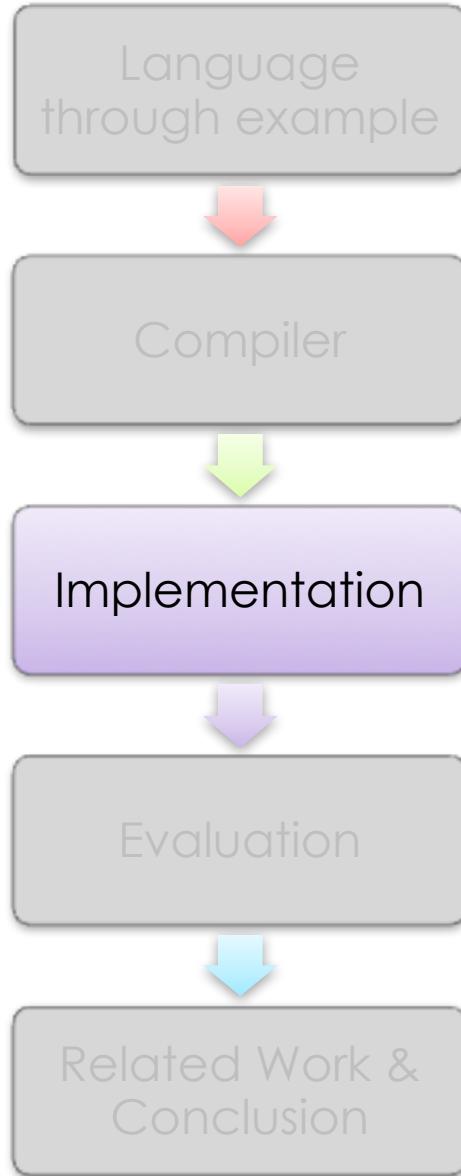


Putting It All Together



Putting It All Together

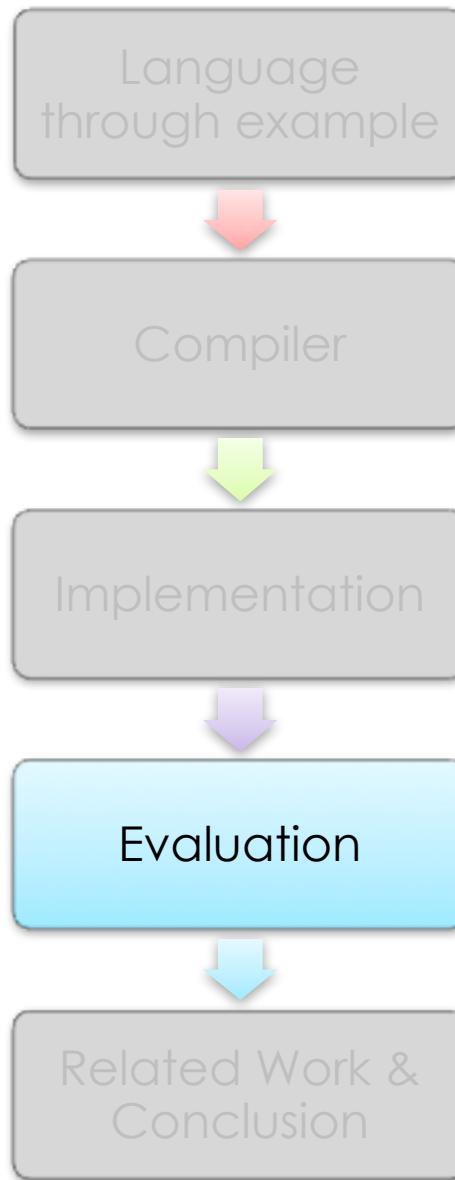




SNAP Implementation

- Compiler written in **Python**
- MILP solver: **Gurobi Optimizer**
- Resulting switch code **NetASM**
(language + software switch)

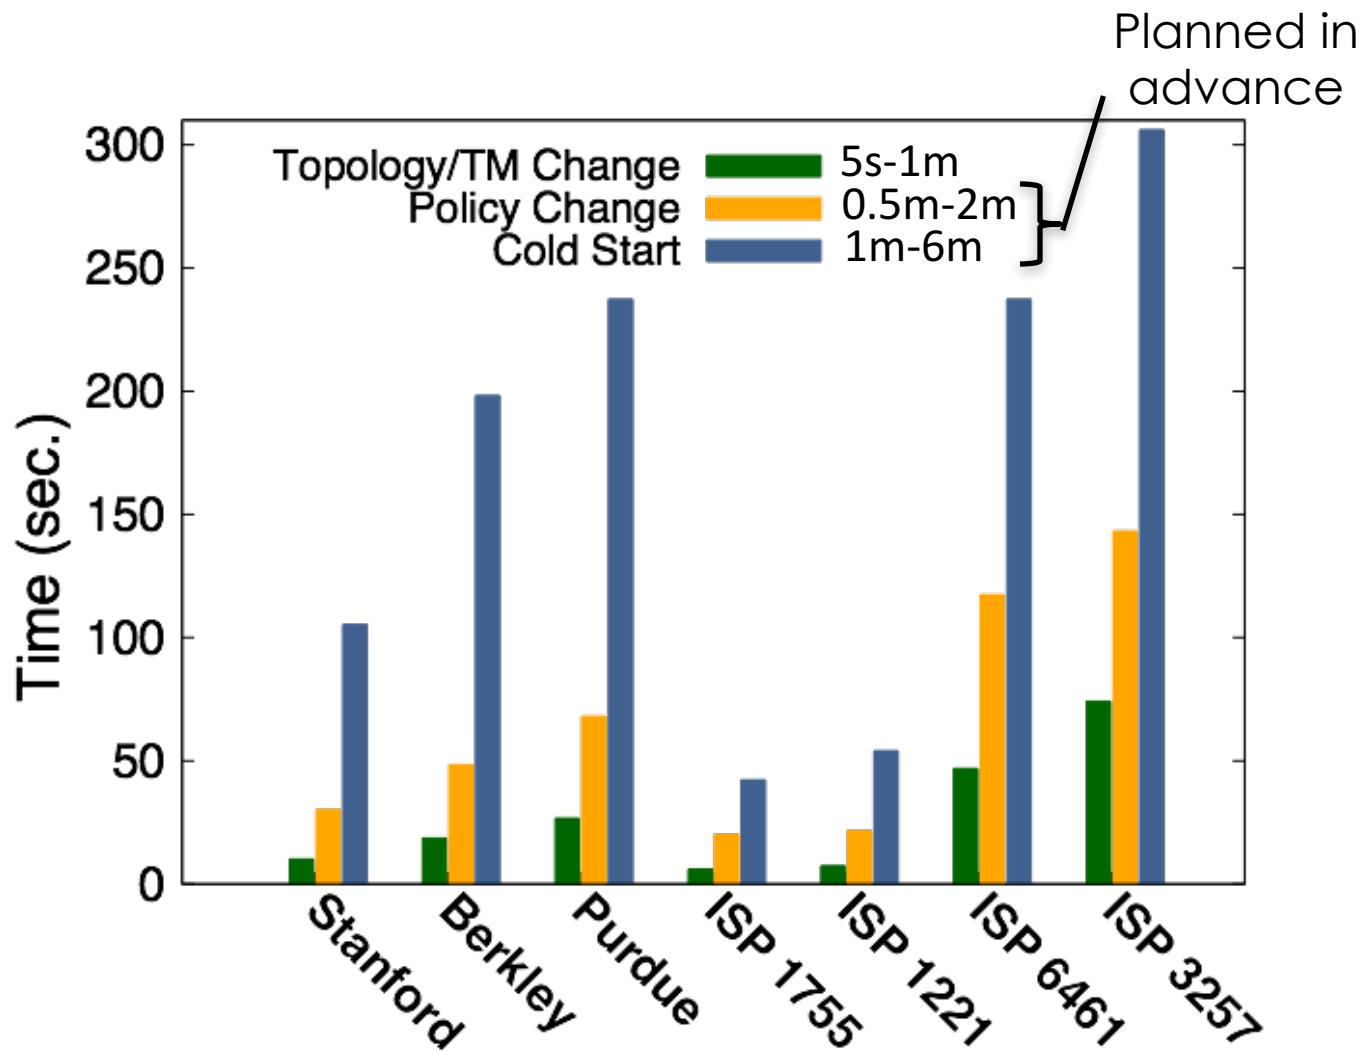
Talk Outline



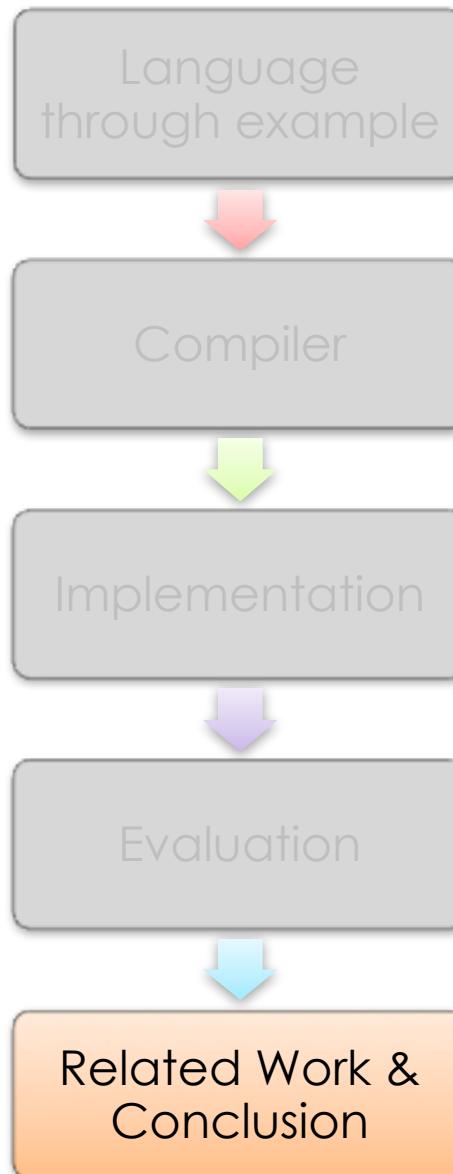
Compiler Evaluation

- 7 campus and ISP topologies
- Order of 100s of switches and links
- Scenarios
 - Cold start (freq. weeks)
 - Policy change (freq. days)
 - Topology/TM change (freq. minutes)

Compiler Evaluation - Results



Talk Outline



Related Work

	Programmable State	Network-Wide	Data Plane State	Joint Placement and Routing	
SNAP (SIGCOMM'16)	✓	✓	✓	✓	Stateful languages
Stateful NetKAT (PLDI'16)	✓	✓	✓	-	Switch level mechanisms
Kinetic (NSDI'15)	✓	✓	-	-	
Domino (SIGCOMM'16)	✓	-	✓	-	
OpenState (SIGCOMM-CCR'14)	✓	-	✓	-	
FAST (HotSDN'14)	✓	-	✓	-	
Merlin (CONECT'14)	-	✓	✓	-	Optimizing placement & routing
Slick (SOSR'15)	✓	✓	-	-	
Stratos (TR'13)	-	✓	✓	-	

Conclusion - SNAP

- A new modular stateful SDN programming language with:
 - One-big switch programming model
 - Persistent global arrays
- Compiler implements algorithms that:
 - Jointly optimize routing and state placement
 - Use efficient IR based on FDDs
- Evaluated about 20 applications